

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import roc_auc_score
import statsmodels.api as sm
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score, cross_validate, train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import precision_recall_curve, average_precision_score
%matplotlib inline
```

```

In [2]: # Считает джини фактора или предсказаний
def gini(y_true, y_score):
    roc_score = roc_auc_score(y_true, y_score)
    return roc_score*2 - 1

def gini_scorrer(clf, X, y):
    y_score = [p[1] for p in clf.predict_proba(X)]
    return gini(y, y_score)

def get_factors_gini(data, target):
    factors_gini = list()
    for c in data.columns:
        g = gini(data[target], data[c])
        factors_gini.append([c,g])
    factors_gini = sorted(factors_gini, key = lambda x: x[1])
    return factors_gini

def get_base_line_score(data, target):
    rf = RandomForestClassifier()
    cv_score = cross_val_score(rf, X = data.drop([target], axis=1), y=data[target], scoring=gini_scorrer, cv=1)
    return cv_score

def get_feature_importance(data, target):
    rf = RandomForestClassifier()
    rf.fit(X = data.drop([target], axis=1), y= data[target])
    return sorted([v,i] for v,i in zip(data.drop([target], axis=1).columns, rf.feature_importances_)), key=lambda x: x[1]

def show_corr(data, th=None):
    if th:
        sns.heatmap(data.corr()>th)
    else:
        sns.heatmap(data.corr())

# Вычисление VIF
def vif(data):
    """
    Calculate VIFs
    :param data: Data
    :return: List of VIFs for each variable in format: Variable - VIF
    """
    res = list()
    for i in range(data.shape[1]):
        res.append([data.columns[i], variance_inflation_factor(data.as_matrix(), i)])
    return res

def rocs(y_true, y_pred, colors, names, name=''):
    """ Plot several ROC curves

    Keyword arguments:
    y_true -- array of true values arrays
    y_pred -- array of predicted values arrays
    colors -- array of colors for curves
    names -- array of names for curves
    name -- name of graph (default empty string)
    """
    if (len(y_true) != len(y_pred)) or (len(y_true) != len(colors)):
        print(len(y_true), len(y_pred), len(colors))
        raise BaseException
    plt.figure(figsize=[10, 10])
    ginis = []

    for i in range(0, len(y_true)):
        fpr, tpr, _ = roc_curve(y_true=y_true[i], y_score=y_pred[i], pos_label=None)
        gini_value = gini(y_true=y_true[i], y_score=y_pred[i])
        ginis.append(gini_value)
        plt.plot(fpr, tpr, color=colors[i], label=names[i] + ' (Gini = %0.4f)' % gini_value)

    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic ' + name)
    plt.legend(loc="lower right")
    plt.show()

def pr_curves(y_true, y_pred, colors, names, name=''):
    """ Plot several PR curves

    Keyword arguments:
    y_true -- array of true values arrays
    y_pred -- array of predicted values arrays
    colors -- array of colors for curves
    names -- array of names for curves
    name -- name of graph (default empty string)
    """

```

```
plt.figure(figsize=[10, 10])

for i in range(0, len(y_true)):
    precision, recall, thresholds = precision_recall_curve(y_true=y_true[i], probas_pred=y_pred[i])
    avg = average_precision_score(y_true[i], y_pred[i])
    plt.plot(recall, precision, lw=1.5, color=colors[i], label=names[i] + ' AUC={0:0.4f}'.format(avg))
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.legend(loc="upper right")
plt.show()
```

```
In [3]: rename_dict = {0: 'prescreaning',
1: 'MA1', 2: 'MA2', 3: 'MA3', 4: 'MA4', 5: 'MA5', 6: 'MA6',
7: 'MAN1', 8: 'MAN2', 9: 'MAN3', 10: 'MAN4', 11: 'MAN5', 12: 'MAN6', 13: 'MAN7', 14: 'MAN8',
15: 'ED', 16: 'D', 17: 'AMFM', 18: 'label'
}
```

```
In [4]: retinopathy = pd.read_excel('retinopathy.xlsx', sheet_name=1, header=None)
data = retinopathy.rename(rename_dict, axis=1)
```

```
In [5]: len(data.columns)
```

```
Out[5]: 19
```

## Base line

```
In [6]: get_base_line_score(data, 'label')
```

```
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
Out[6]: array([0.42545733, 0.43578158, 0.41600708])
```

```
In [7]: get_feature_importance(data, 'label')
```

```
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

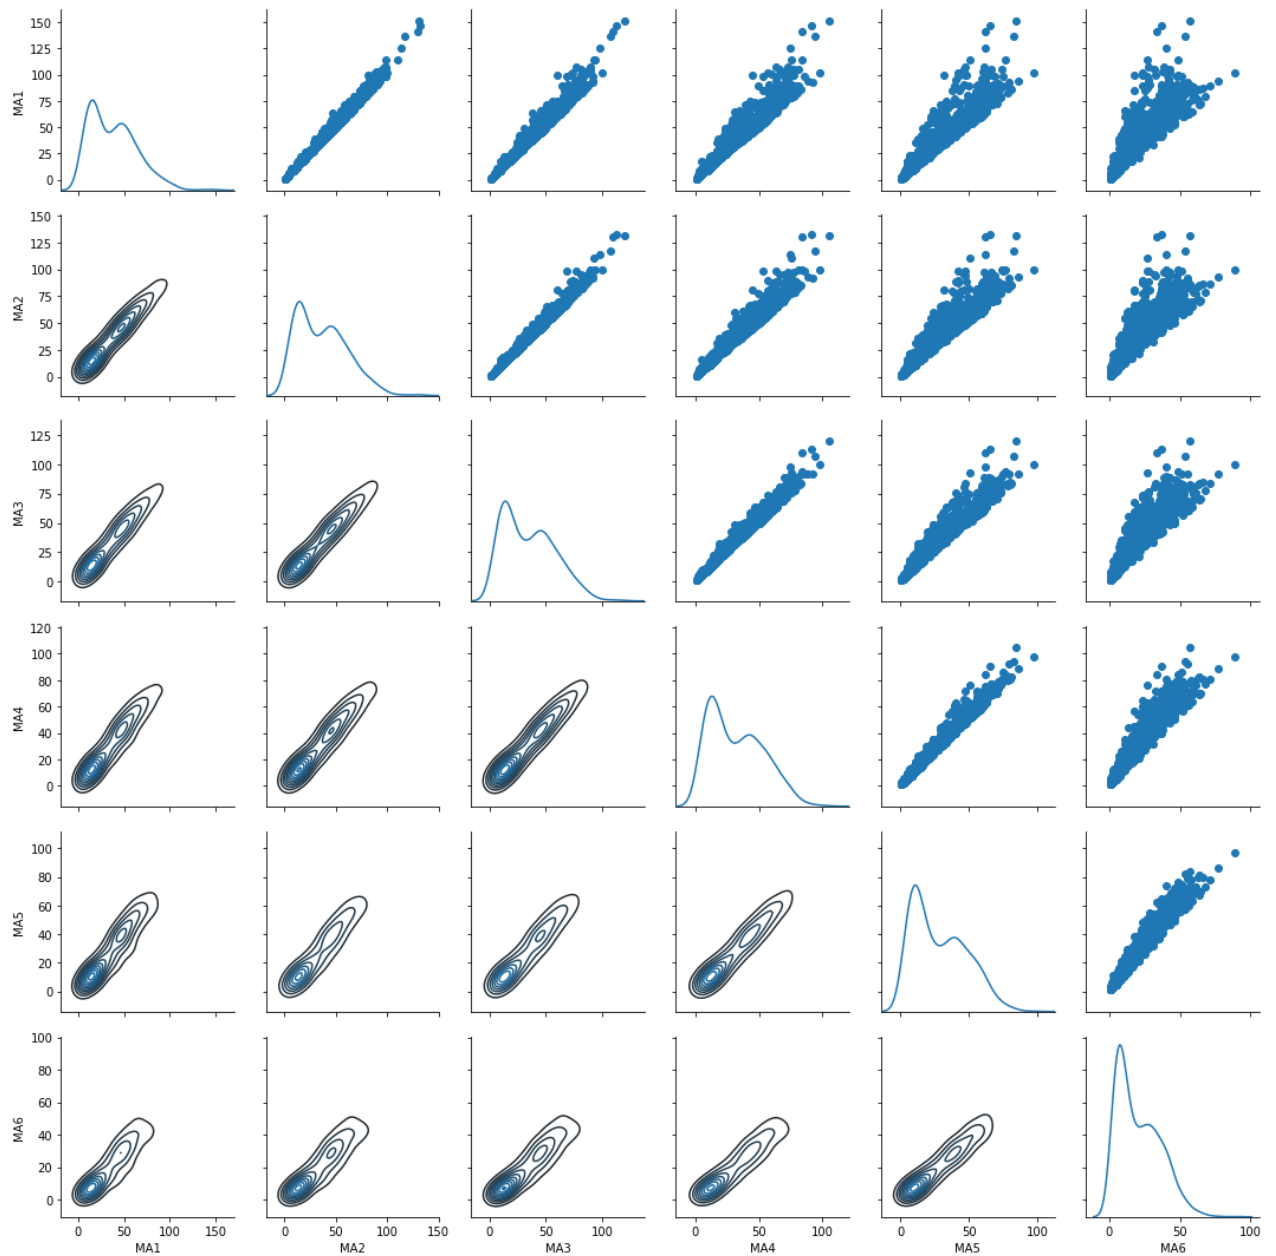
```
Out[7]: [['MA1', 0.09779023323734454],
['MAN1', 0.09035458154793066],
['MAN3', 0.07979423015860702],
['ED', 0.0719682894638727],
['MAN2', 0.06931311572511288],
['D', 0.06383256044214967],
['MAN4', 0.062322423193816597],
['MAN7', 0.06179950175953459],
['MA3', 0.05693749767958459],
['MA2', 0.05652083100143793],
['MAN5', 0.05445902180296812],
['MA6', 0.053522623823726884],
['MA5', 0.04684100424418835],
['MA4', 0.04508688549940245],
['MAN6', 0.04391757534699079],
['MAN8', 0.03131866020755311],
['AMFM', 0.008157684571393584],
['prescreaning', 0.006063280294385651]]
```

## Анализ переменных

### Распределение

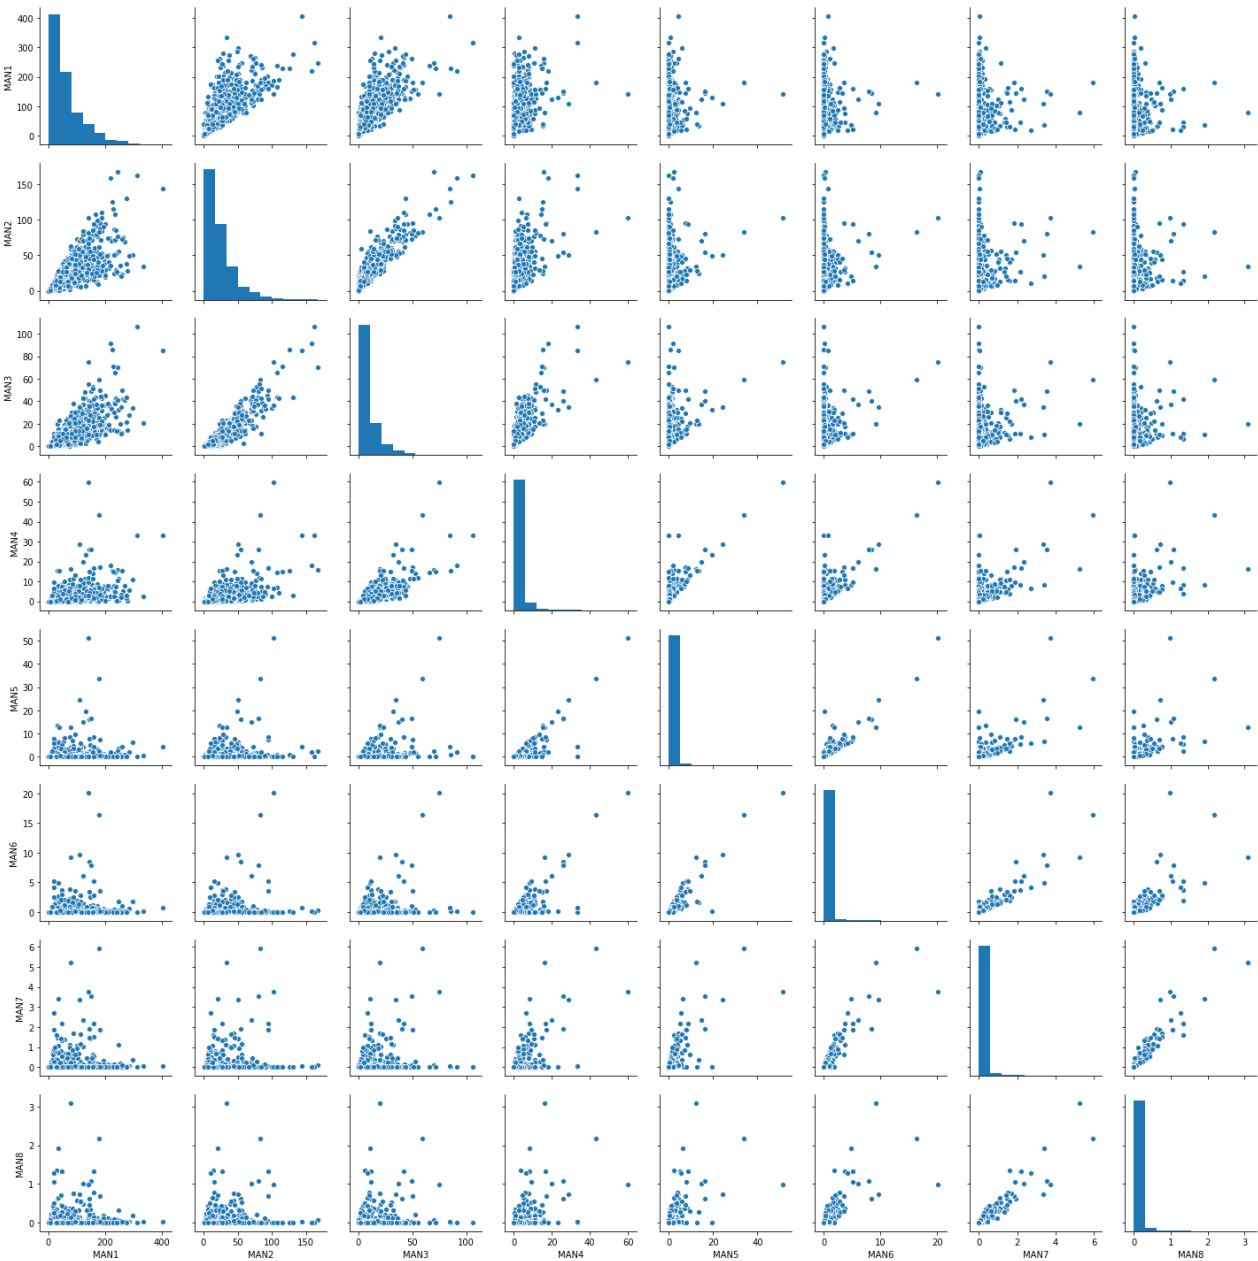
```
In [8]: g = sns.PairGrid(data[['MA1', 'MA2', 'MA3', 'MA4', 'MA5', 'MA6']])
g.map_diag(sns.kdeplot)
g.map_upper(plt.scatter)
g.map_lower(sns.kdeplot)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x1750804c9c8>
```




```
In [9]: sns.pairplot(data[['MAN1', 'MAN2', 'MAN3', 'MAN4', 'MAN5', 'MAN6', 'MAN7', 'MAN8']])
```

Out[9]: <seaborn.axisgrid.PairGrid at 0x1750c3b6b48>



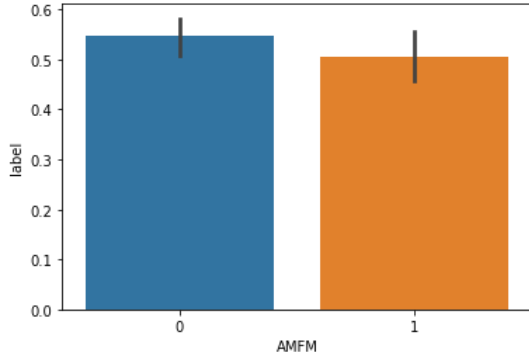
```
In [10]: g2 = sns.PairGrid(data[['D', 'ED']])
g2.map_diag(sns.kdeplot)
g2.map_upper(plt.scatter)
g2.map_lower(sns.kdeplot)
```

Out[10]: <seaborn.axisgrid.PairGrid at 0x1750e3366c8>




```
In [11]: # Бинарные переменные
sns.barplot(x='AMFM', y='label', data=data)
```

Out[11]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1750e2f7888>



```
In [12]: sns.barplot(x='prescreening', y='label', data=data)
```

Out[12]: <matplotlib.axes.\_subplots.AxesSubplot at 0x175102ef408>





```
In [13]: original_features = data.columns.values[:-1]
```

## Анализ корреляций

```
In [14]: corrs = data.corr()
plt.figure(figsize=[10,10])
sns.heatmap(corrs)

plt.figure(figsize=[10,10])
sns.heatmap(corrs>0.7)
```

Out[14]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1751049ef88>



## Новые фичи

```
In [15]: #хочется логарифмировать всё кроме бинарных признаков и ED, D
to_be_logged = ['MA1', 'MA2', 'MA3', 'MA4', 'MA5', 'MA6', 'MAN1',
               'MAN2', 'MAN3', 'MAN4', 'MAN5', 'MAN6', 'MAN7', 'MAN8']
logged = ['{}_log'.format(c) for c in to_be_logged]

for c in to_be_logged:
    data['{}_log'.format(c)] = np.log1p(data[c])
```

```
In [16]: # и попробуем ещё все отношения
from itertools import combinations

feature_comb = list(combinations(original_features,2))

for fc in feature_comb:
    data['{}/{}'.format(fc[0],fc[1])] = data[fc[0]]/data[fc[1]]
    #data['{}/{}'.format(fc[1],fc[0])] = data[fc[1]]/data[fc[0]]
```

```
In [17]: # и произведения
for fc in feature_comb:
    data['{}*{}'.format(fc[0],fc[1])] = data[fc[0]]*data[fc[1]]
    #data['{}/{}'.format(fc[1],fc[0])] = data[fc[1]]/data[fc[0]]
```

```
In [18]: data.replace(np.inf, np.nan, inplace=True)
clear_data = data.dropna(axis=1)
```

## Baseline с новыми фичами

```
In [19]: get_base_line_score(clear_data, 'label')
```

```
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

```
Out[19]: array([0.58771497, 0.67740169, 0.60372502])
```

```
In [20]: all_features_importance = get_feature_importance(clear_data, 'label')
all_features_importance
```

```
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

## Однофакторный анализ

```
In [21]: new_features_gini = get_factors_gini(clear_data, 'label')
new_features_gini_df = pd.DataFrame(new_features_gini, columns=['variable', 'gini'])
new_features_gini_df['gini_abs'] = new_features_gini_df['gini'].abs()
new_features_gini_df.sort_values('gini_abs', ascending=False, inplace=True)
```



```
In [22]: new_features_gini_df.iloc[:40]
```

```
Out[22]:
```

	variable	gini	gini_abs
252	label	1.000000	1.000000
251	MA1/MA3	0.577363	0.577363
250	MA1/MA4	0.559989	0.559989
249	MA1/MA2	0.520602	0.520602
248	MA2/MA4	0.467615	0.467615
247	MA1/MA5	0.454296	0.454296
246	MA2/MA3	0.437529	0.437529
245	MA2/MA5	0.382924	0.382924
244	MA3/MA4	0.377559	0.377559
243	MA1/MA6	0.376933	0.376933
0	prescreaning/MA1	-0.332642	0.332642
242	MA1_log	0.324602	0.324602
241	MA1	0.324602	0.324602
240	MA1*ED	0.323607	0.323607
239	MA2/MA6	0.323595	0.323595
238	MA1/D	0.323033	0.323033
237	MA1/ED	0.322538	0.322538
236	MA3/MA5	0.317552	0.317552
235	MA1*MA2	0.308135	0.308135
1	prescreaning/MA2	-0.307442	0.307442
234	MA1*D	0.301024	0.301024
233	MA2_log	0.294920	0.294920
232	MA2	0.294920	0.294920
231	MA2/D	0.294715	0.294715
230	MA2*ED	0.293952	0.293952
229	MA2/ED	0.293750	0.293750
228	MA1*MA3	0.291335	0.291335
2	prescreaning/MA3	-0.278984	0.278984
227	MA3/MA6	0.278626	0.278626
226	MA2*MA3	0.276880	0.276880
225	MA2*D	0.273255	0.273255
224	MA1*MA4	0.272904	0.272904
222	MA3	0.262012	0.262012
223	MA3_log	0.262012	0.262012
221	MA3/D	0.261799	0.261799
220	MA3*ED	0.261567	0.261567
219	MA3/ED	0.261432	0.261432
218	MA2*MA4	0.257960	0.257960
217	MA1*MA5	0.255857	0.255857
216	prescreaning*MA1	0.249099	0.249099

## значимость

```
In [23]: all_features = set(data.columns) - {'label'}
```

```
In [24]: import statsmodels.api as sm
import statsmodels.formula.api as smf

one_dim_models_res = []
for feature_name in all_features:
    model = smf.logit(f'label ~ {feature_name}', data).fit()
    model.summary()
    one_dim_models_res.append({
        'feature': feature_name,
        'AIC': model.aic,
        'LLR': model.llr,
        'LLR_pval': model.llr_pvalue
    })

one_dim_models_res = pd.DataFrame(one_dim_models_res).set_index('feature')
```

```
Iteration 9
Optimization terminated successfully.
Current function value: 0.654189
Iterations 5
Optimization terminated successfully.
Current function value: 0.666627
Iterations 8
Optimization terminated successfully.
Current function value: 0.689542
Iterations 4
Optimization terminated successfully.
Current function value: 0.646857
Iterations 11
Optimization terminated successfully.
Current function value: 0.684393
Iterations 5
Optimization terminated successfully.
Current function value: 0.690939
Iterations 3
Optimization terminated successfully.
```

```
In [25]: def get_significance(pval):
    if pval < 0.001:
        return '***'
    elif pval < 0.01:
        return '**'
    elif pval < 0.05:
        return '*'
    elif pval < 0.1:
        return '.'
    else:
        return ''

one_dim_models_res['significance'] = one_dim_models_res['LLR_pval'].apply(get_significance)
```

```
In [26]: one_dim_models_res[one_dim_models_res['significance'] == '***'].shape
```

```
Out[26]: (282, 4)
```

```
In [27]: one_dim_models_res.loc[one_dim_models_res['significance'] == '***'].sort_values(by='LLR_pval').head(40)
```

```
Out[27]:
```

	AIC	LLR	LLR_pval	significance
feature				
MA1*MA3	1178.898954	414.273080	1.792138e-89	***
MA1*MA4	1218.432787	374.739247	6.551871e-81	***
MA1*MA2	1248.262139	344.909895	1.887185e-74	***
MA1*MA5	1281.724492	311.447542	3.311661e-67	***
MA2*MA4	1309.669649	283.502385	3.698124e-61	***
MA2*MA3	1311.053141	282.118893	7.367977e-61	***
MA2*MA5	1346.056044	247.115991	2.751578e-53	***
MA1*MA6	1365.792001	227.380033	5.096439e-49	***
MA3*MA4	1395.630717	197.541317	1.433504e-42	***
MA3*MA5	1403.752421	189.419613	8.146689e-41	***
MA2*MA6	1410.371585	182.800449	2.191150e-39	***
MA1/MAN7	1417.604906	173.567128	2.043511e-38	***
MA1*MAN7	1415.359521	177.812513	2.617259e-38	***
MA1/MAN8	1422.994089	168.177945	3.024289e-37	***
MA1*MAN8	1423.591901	169.580133	1.567864e-36	***
MA2*MAN7	1427.637325	165.534709	1.171061e-35	***
MA2/MAN7	1430.451260	160.720774	1.258715e-35	***
MA2/MAN8	1436.474537	154.697497	2.557793e-34	***
MA2*MAN8	1436.639633	156.532401	1.026624e-33	***
MA3*MAN7	1439.191995	153.980040	3.648589e-33	***
MA3/MAN7	1443.803286	147.368748	9.983503e-33	***
MA3/MAN8	1450.115605	141.056429	2.344151e-31	***
MA1/MAN6	1450.189392	140.982642	2.432251e-31	***
MA1*MAN6	1447.791624	145.380410	2.613336e-31	***
MA1*AMFM	1449.054817	144.117217	4.893557e-31	***
MA3*MAN8	1449.207454	143.964580	5.278888e-31	***
MA4*MAN7	1450.201273	142.970762	8.646964e-31	***
MA1/MA6	1453.510784	137.661251	1.280086e-30	***
MA3*MA6	1451.815122	141.356912	1.926973e-30	***
MA1/MA5	1454.522682	136.649352	2.123099e-30	***
MA4/MAN7	1456.321933	134.850101	5.220024e-30	***
MA5*MAN7	1457.101642	136.070392	2.658697e-29	***
MA1/AMFM	1460.130416	131.041618	3.504888e-29	***
MA4/MAN8	1463.242512	127.929522	1.661335e-28	***
MA1/MAN5	1463.446052	127.725982	1.839312e-28	***
MA4*MAN8	1461.218378	131.953656	2.051319e-28	***
MA1/MAN1	1463.677605	127.494429	2.065078e-28	***
MA2*MAN6	1461.385825	131.786209	2.229061e-28	***
MA6*MAN7	1462.225552	130.946482	3.381420e-28	***
MA1*MAN5	1462.546050	130.625984	3.964333e-28	***

## Регрессия

```
In [28]: def test_model(train, test, model):
    rocs(y_true = [train['label'], test['label']],
        y_pred = [fitted_model.predict(sm.add_constant(train[predictors])), fitted_model.predict(sm.add_constant(
            colors = ['blue', 'red'],
            names = ['train', 'test'], name='')
        print('Train gini:', gini(train['label'], fitted_model.predict(sm.add_constant(train[predictors]))))
        print('Test gini:', gini(test['label'], fitted_model.predict(sm.add_constant(test[predictors]))))
        print(fitted_model.summary())
        show_corr(clear_data[predictors])
        vif_df = pd.DataFrame(vif(clear_data[predictors]), columns=['variable', 'vif'])
        print(vif_df.head(10))
```


```
In [29]: train, test = train_test_split(clear_data, test_size=0.3, stratify = clear_data['label'], random_state=42)
```

```
In [30]: predictors = ['MA1/MA3', 'AMFM']
#predictors = ['MA1/MA4', 'MA1*ED']

model = sm.Logit(train['label'], sm.add_constant(train[predictors]))
fitted_model = model.fit()
test_model(train, test, fitted_model)
```

c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.  
 return ptp(axis=axis, out=out, \*\*kwargs)

Optimization terminated successfully.  
 Current function value: 0.556833  
 Iterations 7



c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.  
 return ptp(axis=axis, out=out, \*\*kwargs)

c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.  
 return ptp(axis=axis, out=out, \*\*kwargs)

Train gini: 0.5984699453551914  
 Test gini: 0.5651836348906294


#### Logit Regression Results

```
=====
Dep. Variable:          label    No. Observations:          802
Model:                Logit      Df Residuals:            799
Method:                MLE       Df Model:              2
Date:                  Fri, 03 Jan 2020    Pseudo R-squ.:          0.1942
Time:                  13:27:33    Log-Likelihood:         -446.58
converged:              True      LL-Null:               -554.22
Covariance Type:       nonrobust    LLR p-value:            1.795e-47
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-14.9697	1.384	-10.819	0.000	-17.682	-12.258
MA1/MA3	14.2637	1.315	10.848	0.000	11.686	16.841
AMFM	-0.7263	0.183	-3.958	0.000	-1.086	-0.367


```
=====
variable      vif
0  MA1/MA3    1.529027
1    AMFM     1.529027
=====
```

c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\ipykernel\_launcher.py:44: FutureWarning: Method .as\_matrix will be removed in a future version. Use .values instead.



```
In [31]: g3 = sns.PairGrid(clear_data[predictors])
g3.map_diag(sns.kdeplot)
g3.map_upper(plt.scatter)
g3.map_lower(sns.kdeplot)
```

```
Out[31]: <seaborn.axisgrid.PairGrid at 0x17510bbca88>
```




## Влияние выбросов

```
In [32]: from statsmodels.genmod.generalized_linear_model import GLM
from statsmodels.genmod import families

final_features = np.array(train[predictors])
cooks_model = GLM(train['label'], final_features, family=families.Binomial()).fit()

infl = cooks_model.get_influence()
plt.figure(figsize=(12, 4))
sns.scatterplot(fitted_model.fittedvalues, infl.cooks_distance[0])
```


```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x17510ac68c8>
```



```
In [33]: new_train = train[infl.cooks_distance[0] < 0.01]
print('Fitting on {} / {} of samples'.format(len(new_train), len(train)))
model_cooks = sm.Logit(new_train['label'], sm.add_constant(new_train[predictors])).fit()
test_model(new_train, test, model_cooks)
```

Fitting on 802 / 802 of samples  
 Optimization terminated successfully.  
 Current function value: 0.556833  
 Iterations 7

c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.  
 return ptp(axis=axis, out=out, \*\*kwargs)



c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.  
 return ptp(axis=axis, out=out, \*\*kwargs)  
 c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.  
 return ptp(axis=axis, out=out, \*\*kwargs)  
 c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\ipykernel\_launcher.py:44: FutureWarning: Method .as\_matrix will be removed in a future version. Use .values instead.

Train gini: 0.5984699453551914  
 Test gini: 0.5651836348906294


#### Logit Regression Results

```
=====
Dep. Variable:          label    No. Observations:          802
Model:                Logit      Df Residuals:            799
Method:               MLE        Df Model:              2
Date:                Fri, 03 Jan 2020    Pseudo R-squ.:        0.1942
Time:                13:27:38    Log-Likelihood:       -446.58
Converged:            True        LL-Null:             -554.22
Covariance Type:      nonrobust    LLR p-value:         1.795e-47
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-14.9697	1.384	-10.819	0.000	-17.682	-12.258

MA1/MA3	14.2637	1.315	10.848	0.000	11.686	16.841
AMFM	-0.7263	0.183	-3.958	0.000	-1.086	-0.367

```
=====
variable      vif
0  MA1/MA3    1.529027
1    AMFM     1.529027
```



Выбросы не оказывают большого влияния, удалять их не будем


## Оценка качества модели

```
In [34]: test_scores = fitted_model.predict(sm.add_constant(test[predictors]))
train_scores = fitted_model.predict(sm.add_constant(train[predictors]))
```

```
c:\users\xiaomi\appdata\local\programs\python\python37\lib\site-packages\numpy\core\fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
```

## Разделяющая способность


```
In [35]: rocs(y_true=[test['label'], train['label']], y_pred=[test_scores, train_scores],
           colors=['red', 'blue'], names=['test', 'train'])
```





## Средняя точность (PR-кривая)

```
In [36]: pr_curves(y_true=[test['label'], train['label']], y_pred=[test_scores, train_scores],
                  colors=['red', 'blue'], names=['test', 'train'])
```




## Чувствительность и специфичность

```
In [37]: n1 = test['label'].sum()
n = len(test['label'])

sensitivity = np.zeros_like(test_scores)
specificity = np.zeros_like(test_scores)
for j, u in enumerate(sorted(test_scores)):
    sensitivity[j] = np.sum((test_scores >= u) * test['label']) / float(n1)
    specificity[j] = np.sum((test_scores <= u) * (1 - test['label'])) / float(n - n1)

plt.figure(figsize=(6, 6))
sns.lineplot(sorted(test_scores), specificity, label='specificity',)
sns.lineplot(sorted(test_scores), sensitivity, label='sensitivity')
plt.xlabel('Cutoffs')
```

Out[37]: Text(0.5, 0, 'Cutoffs')



## Порог отсека

```
In [38]: sorted_scores = sorted(test_scores)
idx = np.nonzero(np.array(sorted_scores) >= 0.45)[0][0]
print('For cutoff {:.4f}: specificity={:.4f}, sensitivity={:.5f}'.format(
    sorted_scores[idx], specificity[idx], sensitivity[idx]
))
```

For cutoff 0.4522: specificity=0.6832, sensitivity=0.72826

## Выводы и интерпретация

```
In [39]: confint_results = np.exp(pd.concat((fitted_model.params, fitted_model.conf_int()), axis=1))
confint_results.columns = ['coef', 'confint_2.5', 'confint_97.5']
confint_results
```

Out [39]:

	coef	confint_2.5	confint_97.5
<b>const</b>	3.153142e-07	2.094216e-08	4.747507e-06
<b>MA1/MA3</b>	1.565492e+06	1.189549e+05	2.060247e+07
<b>AMFM</b>	4.836963e-01	3.375911e-01	6.930339e-01

### В логистической регрессии

Вероятность наступления события зависит от  $z$ :  $p = \frac{1}{1+e^{-z}}$ , где  $z = a_0 + a_1 * x_1 + \dots + a_n * x_n$

Шанс наступления события зависит от  $z$ :  $\frac{p}{1-p} = e^z = e^{a_0 + a_1 * x_1 + \dots + a_n * x_n}$

Таким образом, коэффициент регрессии  $a_i$  показывает, что при увеличении значения объясняющей переменной на 1, шанс наступления события увеличивается в  $e^{a_i}$  раз

### Интерпретация модели

Разберём значения коэффициентов в полученной модели. Переменные вошедшие в модель:

- MA1/MA4 - отношение количества МА, найденных на уровне p-value 0.5, к количеству МА, найденных на уровне p-value 0.7
- MA1\*ED - результат AM/FM классификации

Видим, что:

- увеличение MA1/MA3 на 1 приводит к увеличению шанса возникновения целевого события в  $e^{14.2637} = 1565475.23$  раз
- положительный (1) результат AM/FM классификации приводит к уменьшению шанса возникновения целевого события в  $e^{0.7263} = 2.07$  раз